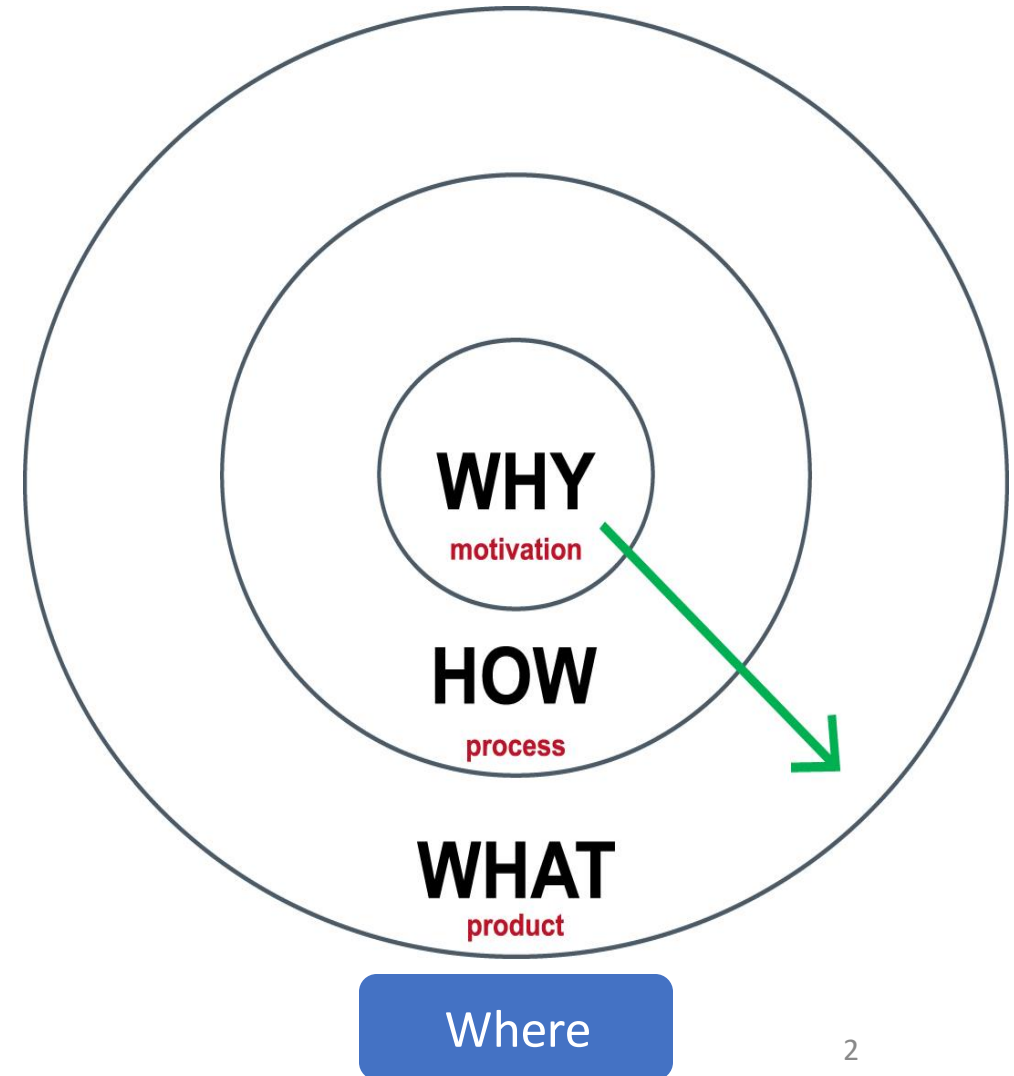


Prompt compression: Why, How, What and Where

Petr Lorenc

Agenda - Start with **WHY** ...

- **WHY** we did it?
- **HOW** was it done?
- **WHAT** is there beneficial for you?
- **WHERE** can I try it?



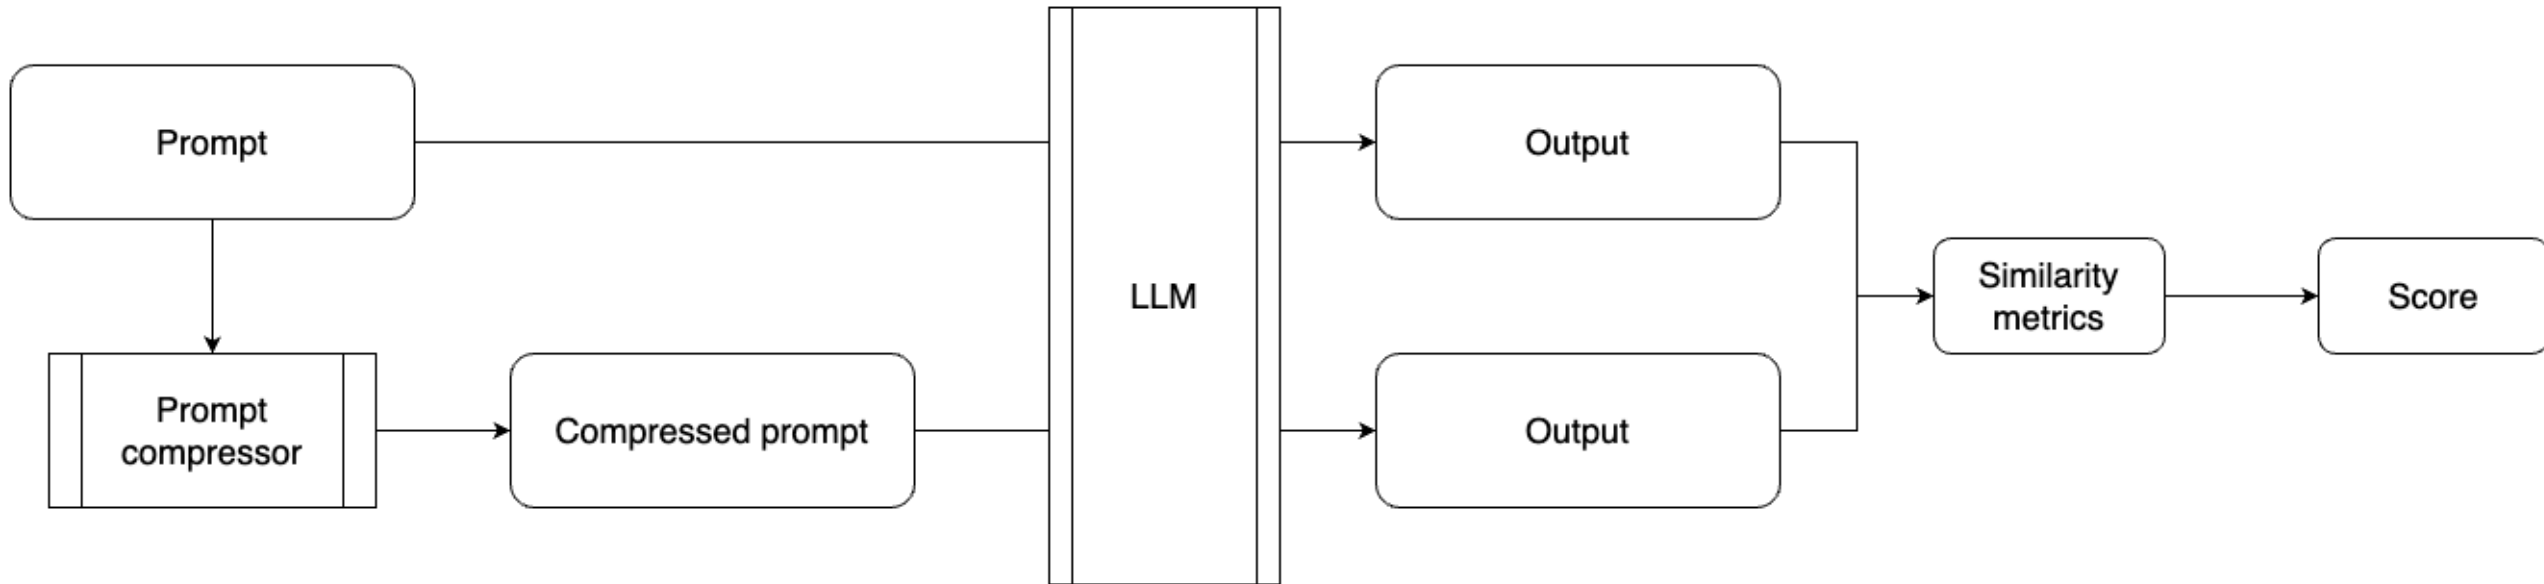
But **Why** is it important for **you**?

- Because we want to reduce a cost
 - It can reduce the token count = reduce the cost for the company
- Because we do not want to wait
 - It can speed up your prompt (less tokens for input)



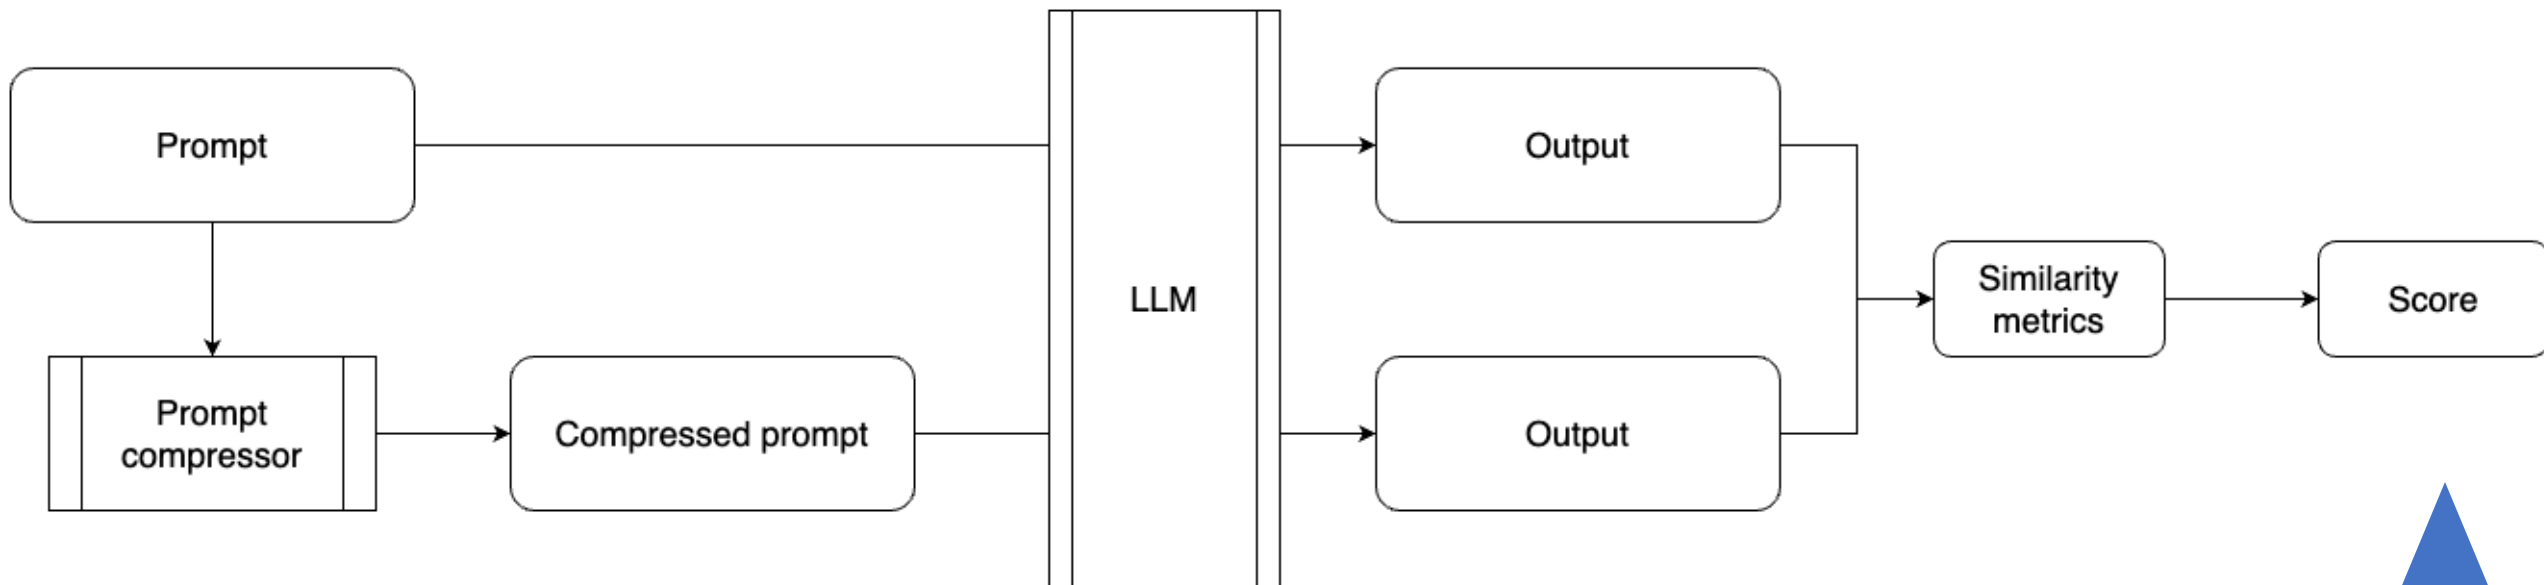
So **How** it was measured?

- We want the same output as without any prompt compression



Let's start from the end ...

- We want the same output as without any prompt compression

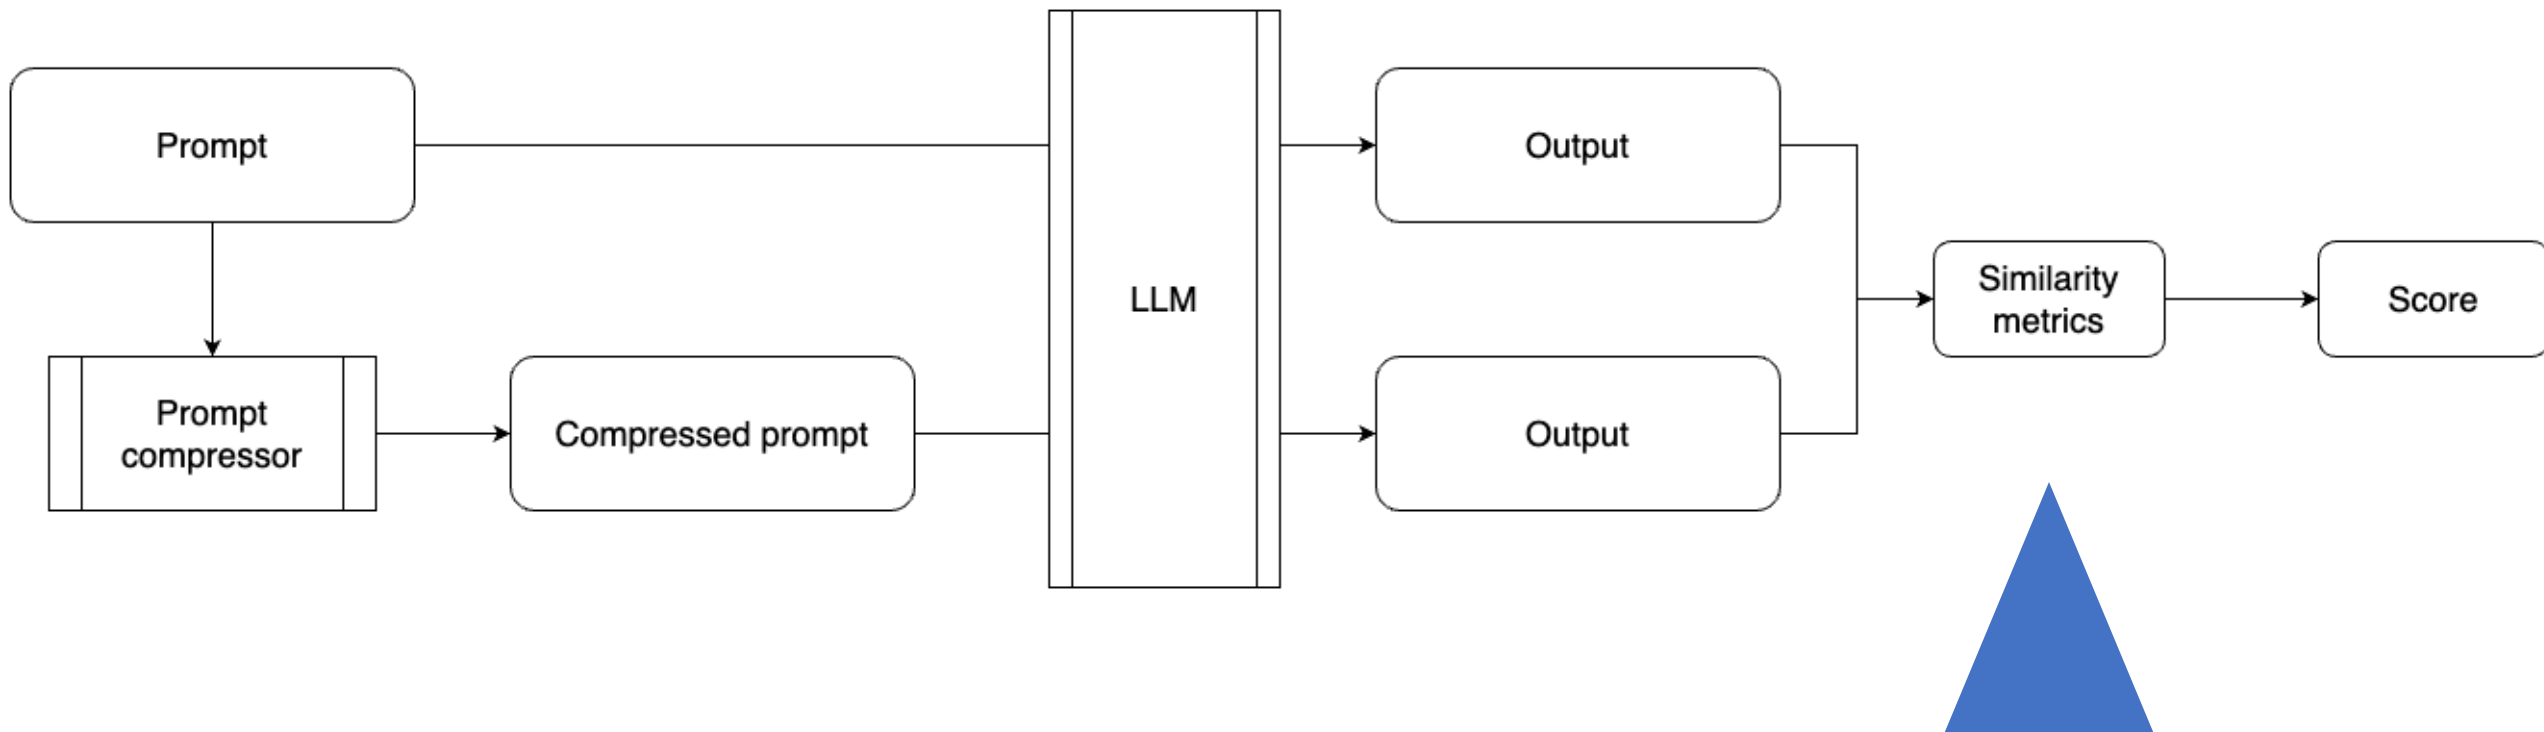


To get **Score** we need some test-data

- **Summarisation (similar to non-RAG applications):**
 - Given a context, the task is to generate a summary that captures the main points of the document. This task aims to evaluate how compression affects the overall understanding of models on the input contexts.
 - <https://huggingface.co/datasets/EdinburghNLP/xsum>
- **QA over Documents (similar to RAG applications):**
 - Given a document, the task is to answer given question.
 - We skip the Retrieval part (not a factor here)
 - <https://huggingface.co/datasets/rag-datasets/rag-mini-bioasq>


We have the dataset, now we need **metrics**

- We want the same output as without any prompt compression



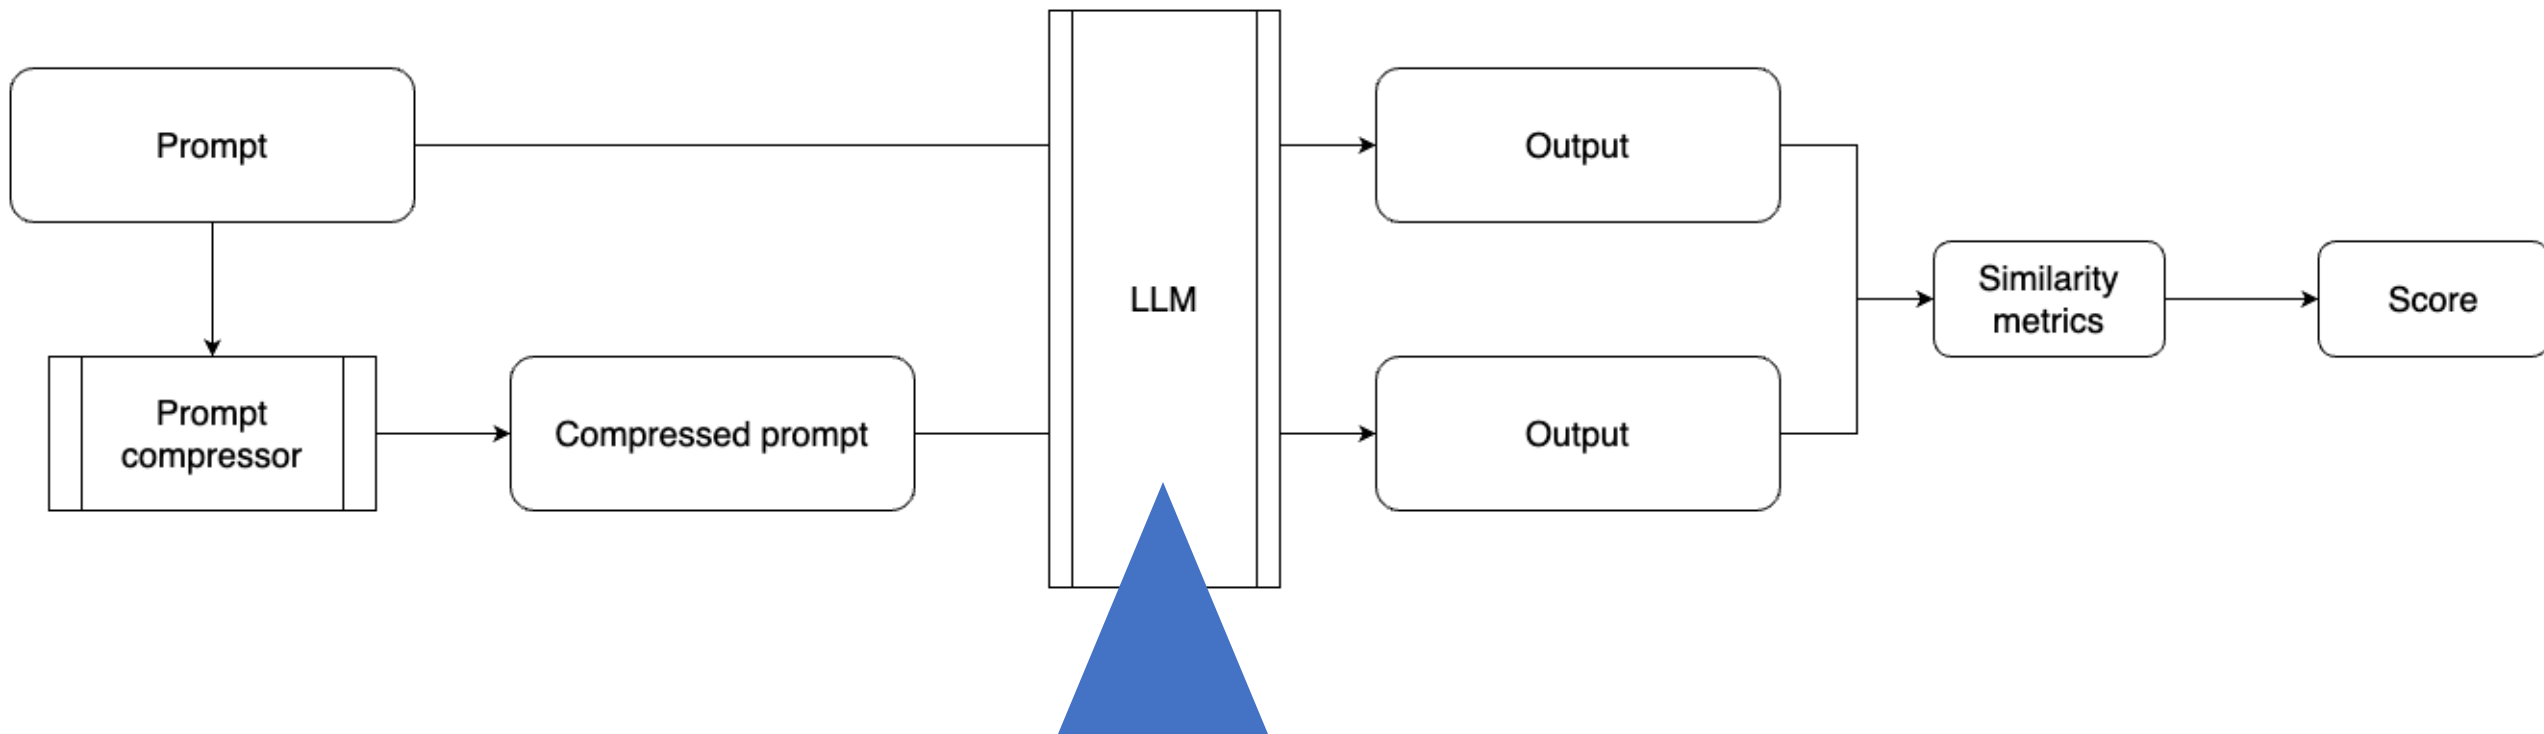
Selection of similarity metrics

- **BM25**
 - Based on tokens overlap
- **GZIP + NCD (Normalized Compression Distance)**
 - Based on difference between the compression rate of concatenated text and each part separately
- **ROUGE-L: Longest Common Subsequence (LCS)**
 - Based on tokens overlap
- **BERTscore F1**
 - Based on semantic similarity overlap of tokens (embeddings)
- **SentenceTransformers (SBERT) + cosine similarity**
 - Based on semantic overlap
- **OpenAI GPT4**
 - Based on “semantic” overlap + reasoning




apiGPTeal family (Azure-based)

- We want the same output as without any prompt compression



And finally, we got to the compression part

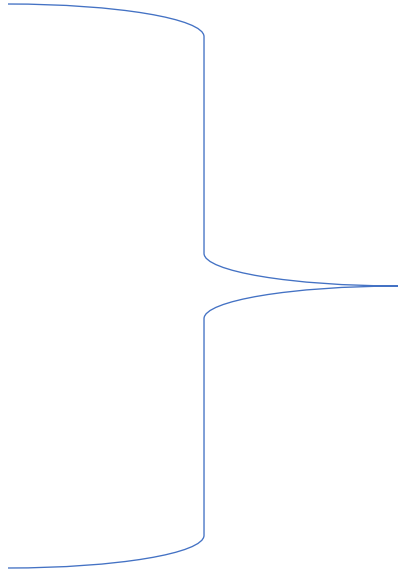
- We want the same output as without any prompt compression



How it was measured?

- **Baselines**
 - Ask GPT4 to compress the prompt
 - Randomly remove 50% (chosen artificially because of other approaches)
- **Used Approaches**
 - LLMingua
 - LongLLMingua
 - LLMingua 2
 - Selective Context
 - Prompt Optimizer
 - GPtrim

How it was measured?

- **Baselines**
 - Ask GPT4 to compress the prompt
 - Randomly remove 50% (chosen artificially because of other approaches)
 - **Used Approaches**
 - LLMingua
 - LongLLMingua
 - LLMingua 2
 - Selective Context
 - Prompt Optimizer
 - Gptrim
- 
- remove additional information


How it was measured?

- **Baselines**

- Ask GPT4 to compress the prompt
- Randomly remove 50% (chosen artificially because of other approaches)

- **Used Approaches**


- LLMlingua
- LongLLMLingua
- LLMlingua 2



Assume well trained LLM model (with perplexity output)
Lower perplexity = lower contribution to the prompt
Lower perplexity = lower surprise of next token
I am hap.. -> ...py VS ...tic

- Selective Context
- Prompt Optimizer

- Gptrim



Remove stop-words (the, a, ...)
Stemming (running -> run)

What was measured?


GPT 4 turbo	GPT 4	Random	LongLLMLingua	LLMLingua	LLMLingua 2	GPTrim
BM25	7.21	5.54	6.41	6.82	11.86	10.9
GZIP	0.53	0.58	0.55	0.54	0.44	0.39
ROUGE	0.53	0.59	0.56	0.559	0.44	0.38
BERTScore F1	0.112	0.127	0.11	0.117	0.08	0.07
SBert	0.100	0.117	0.10	0.116	0.07	0.06
GPT 4 as Judge	0.22	0.246	0.241	0.253	0.14	0.117
Latency	19.69 s	0.0025s	7.0732 s	5.3 s	0.2333 s	0.0096 s
Compression	2690 → 249 (10%)	2690 → 1322 (50%)	2690 → 1223 (45%)	2690 → 828 (30%)	2690 → 1388 (51%)	2690 → 2323 (86%)
HW requirements	None	None	GPU - 17-21G	GPU - 17-21G ml.g5.xlarge	GPU - 1-2GB model	None

What was measured?


GPT 4 turbo	GPT 4	Random	LongLLMLingua	LLMLingua	LLMLingua 2	GPTrim
BM25	7.21	5.54	6.41	6.82	11.86	10.9
GZIP	0.53	0.58	0.55	0.54	0.44	0.39
ROUGE	<div>Is it good result or not?</div>				0.44	0.38
BERTScore F1					0.08	0.07
SBert					0.07	0.06
GPT 4 as Judge					0.14	0.117
Latency					0.2333 s	0.0096 s
Compression	2690 → 249 (10%)	2690 → 1322 (50%)	2690 → 1223 (45%)	2690 → 828 (30%)	2690 → 1388 (51%)	2690 → 2323 (86%)
HW requirements	None	None	GPU - 17-21G	GPU - 17-21G ml.g5.xlarge	GPU - 1-2GB model	None

Is it good result or not?

Baseline



RAG	T=0,0	T=0,2	T=0,4	T=0,6	T=0,8	T=1,0	T=1.2	T=1.4
BM25	16.98	15.69	14.71	14.51	13.90	13.18	12.32	11.69
GZIP	0.23	0.26	0.33	0.42	0.44	0.47	0.53	0.55
ROUGE	0.20	0.23	0.32	0.41	0.45	0.46	0.56	0.58
BERTScore	0.037	0.042	0.05	0.077	0.080	0.089	0.10	0.11
SBert	0.019	0.027	0.045	0.069	0.066	0.062	0.09	0.08
GPT 4	0.068	0.083	0.095	0.14	0.13	0.14	0.17	0.18



RAG		T=0,0	T=0,2	T=0,4	T=0,6	T=0,8	T=1,0	T=1.2	T=1.4
BM25	↑	16.98	15.69	14.71	14.51	13.90	13.18	12.32	11.69
GZIP	↓	0.23	0.26	0.33	0.42	0.44	0.47	0.53	0.55
ROUGE	↓	0.20	0.23	0.32	0.41	0.45	0.46	0.56	0.58
BERTScore	↓	0.037	0.042	0.05	0.077	0.080	0.089	0.10	0.11
SBert	↓	0.019	0.027	0.045	0.069	0.066	0.062	0.09	0.08
GPT 4	↓	0.068	0.083	0.095	0.14	0.13	0.14	0.17	0.18

GPT 4 turbo	LLMLingua 2	GPTrim	T=1,0		GPT 4o	LLMLingua 2	GPTrim	T=1,0	
BM25	11.86	10.9	13.18	↑	BM25	14.27	14.79	13.18	↑
GZIP	0.44	0.39	0.47	↓	GZIP	0.466	0.425	0.47	↓
ROUGE	0.44	0.38	0.46	↓	ROUGE	0.453	0.404	0.46	↓
BERTScore F1	0.08	0.07	0.089	↓	BERTScore F1	0.095	0.0869	0.089	↓
SBert	0.07	0.06	0.062	↓	SBert	0.063	0.0563	0.062	↓
GPT 4 as Judge	0.14	0.117	0.14	↓	GPT 4 as Judge	0.168	0.145	0.14	↓

How it looks like in practise?


Many in **the list have been** found down old mine tunnels **or on** slag heaps where water and **even fire have had the** opportunity to work **up** novel **compounds**. It is another example, the researchers argue, of our **pervasive** influence on the planet.\n



the list have been or on even fire have had the upounds**It** another per planet.

So **WHAT** is there beneficial for you?

- You can save the money spend
 - Company pays for the tokens, so less tokens = less money



- You can reduce the latency of your prompt, but
 - it depends on the method
 - LLMLingua 2 is getting fast but still it is additional model
 - Gptrim has lower performance but it is super fast



- You will get the similar output as having reasonable low temperature



Conclusion - Start with **WHY** ... End with answers

- **WHY** we did it?
 - Because we want work with longer context.
- **HOW** was it done?
 - By comparing output of uncompressed and compressed prompt
- **WHAT** is there beneficial for you?
 - It can save cost and potentially speed up your repeating prompt

